Megha Srivastava

megha@cs.stanford.edu • www.cs.stanford.edu/~megha
735 Campus Dr., Stanford, CA 94305

Education

Stanford University • Stanford, CA — Fall 2020 – *present*
PhD in Computer Science, co-advised by Dorsa Sadigh & Dan Boneh
Supported by the NSF Graduate Research Fellowship (2018 – 2023) and IBM PhD Fellowship (2023-2024)
Current Research Interests: robust machine learning, human-AI interaction, AI for scientific discovery

Massachusetts Institute of Technology • Cambridge, MA — Fall 2023 – Winter 2024
Visiting PhD student at the Computer Science & Artificial Intelligence Laboratory, hosted by Jacob Andreas

Stanford University • Stanford, CA — Fall 2014 – Spring 2019
BS in Computer Science with Honors and Distinction, Minor in Creative Writing
MS in Computer Science (Artificial Intelligence), advised by Percy Liang & Tatsunori Hashimoto
Ben Wegbreit Prize for Best Undergraduate Honors Thesis in Computer Science
Bing Overseas Studies Program at Oxford University (Logic & Computability) in Spring 2017

Research Experience

AI Resident – Microsoft Research — Fall 2019 – Fall 2020
Redmond, WA
- Quantifying effects of nondeterminism in neural network training (Mentors: Besmira Nushi, Ece Kamar, Eric Horvitz).

Research Intern – Google Research — Summer 2019
Los Angeles, CA
- Improving robustness of image understanding models with the Research & Machine Intelligence team.

Research Intern – ETH Zürich Learning & Adaptive Systems Group — Summer 2018
Zürich, Switzerland
- Active learning and adaptive hypothesis testing (Mentors: Andreas Krause, Hoda Heidari)
- Supported by the ETH Zürich Student Summer Research Fellowship.

Research Assistant – Vision & Perception Neuroscience Group — Summer 2016 – Spring 2017
Stanford, CA
- Generalization and perceptual invariances in both human and artificial vision models (Mentor: Kalanit Grill-Spector)
- Supported by the Bio-X Undergraduate Research Award.

Awards & Fellowships

- Human Robot Interaction (HRI) Pioneers, 2025
- IBM PhD Fellowship, 2023
- American Association for the Advancement of Science (AAAS) Mass Media Fellowship Finalist, 2019
- International Conference on Machine Learning (ICML) Best Paper Runner-Up Award, 2018
- National Science Foundation Graduate Research Fellowship, 2018
- Ben Wegbreit Prize for Best Undergraduate Honors Thesis in Computer Science (Stanford), 2018
- Phi Beta Kappa, 2018
- Bio-X Research Award (Stanford), 2016
- Lunsford Award for Oral Presentation of Research Finalist (Stanford), 2016
- Intel Science Talent Search Semifinalist (High School), 2013

Activities & Service

- Science Small Groups Mentor (2024)
- Course Assistant: CS 255 (Introduction to Cryptography, 2025), CS 221 (Introduction to Artificial Intelligence, 2018)
- Stanford AI Lab Blog Editor, 2020-present
- Organizer for Women in Machine Learning (WiML) 2023, co-located with NeurIPS 2023
  Student Program & Travel Funding Chair
- Reviewer for NeurIPS (2020 - curr.), ICLR (2021 - curr.), ICML (2021 - curr.), CoRL 2023, L4DC 2024, CHI 2025
- Stanford Women in Computer Science (Vice President), 2017 – 2018
- Photography showcased at Stanford Art of Science Exhibition, Frost Music Festival, dnjGallery at Bergamot Station, and Stanford UNICEF Arts Gallery

contd.

Preprints

Megha Srivastava, Cedric Colas, Dorsa Sadigh, Jacob Andreas. *"Policy Learning via Language Bottleneck"*, 2025
Spotlight at Training Agents with Foundation Models Workshop (RLC 2024).

Selected Publications (See Google Scholar for full list, * denotes equal contribution)

Megha Srivastava*, Reihaneh Iranmanesh*, Yuchen Cui, Deepak Gopinath, Emily Sumner, Andrew Silva, Laporsha Dees, Guy Rosman, Dorsa Sadigh. *"Shared Autonomy for Proximal Teaching"*, ACM/IEEE International Conference on Human-Robot Interaction (HRI), 2025

Megha Srivastava, Simran Arora, Dan Boneh. *"Optimistic Verifiable Training by Controlling Hardware Nondeterminism"*, Advances in Neural Information Processing Systems 37 (NeurIPS) , 2024.

Neil Perry*, **Megha Srivastava***, Deepak Kumar, Dan Boneh. *"Do Users Write More Insecure Code with AI Assistants?"* ACM Conference on Computer and Communications Security (CCS), 2023 **mlsec.org Top-100 Computer Security Papers**

Megha Srivastava and Noah Goodman. *"Question Generation for Adaptive Education"* Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL), 2021.

Megha Srivastava, Tatsunori Hashimoto, Percy Liang. *"Robustness to Spurious Correlations via Human Annotations"* Proceedings of the 37th International Conference on Machine Learning (ICML), 2020.

Tatsunori Hashimoto, **Megha Srivastava**, Hongseok Namkoong, Percy Liang. *"Fairness Without Demographics in Repeated Loss Minimization."* Proceedings of the 35th International Conference on Machine Learning (ICML), 2018. **Best Paper Runner-Up Award**

Invited Talks

1. "Security of Modern Machine Learning" *AI4All, June 2025*
2. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *Google Privacy in ML Seminar, February 2025*
3. "Security of Modern Machine Learning" *DaVita x Stanford Human-Centered AI, February 2025*
4. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *AI and Cryptography Workshop, Joint Mathematics Meeting, January 2025*
5. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *University of Washington Security Seminar, January 2025*
6. "New Challenges of Trust with Large-Scale AI Systems" *University of Chicago, November 2024*.
7. Beyond Instruction Following: Language and Human-Robot Interaction *Stanford University (Guest Lecture for CS 329X: Human-Centered NLP), October 2024*
8. "Security of Modern Machine Learning" *Stanford Cybersecurity & Privacy Festival, October 2024*
9. "Policy Learning with a Language Bottleneck" ***Spotlight Talk*** *at Training Agents with Foundation Models Workshop (RLC 2024), August 2024*.
10. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *University of Toronto, July 2024*
11. "Security of Modern Machine Learning" *Accenture Responsible AI Course, July 2024*
12. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *Stanford Security Forum, April 2024*
13. "Challenges in Human-AI Interaction for Information-Seeking Tasks" *Rising Stars in Machine Learning Workshop, November 2023*
14. "Do Users Write More Insecure Code with AI Assistants?" *University of Maryland College Park (Guest Lecture for Large Language Models, Security, and Privacy seminar), October 2023*.
15. "Robustness to Spurious Correlations via Human Annotations" *Two Sigma PhD Symposium, June 2023*.

16. "Do Users Write More Insecure Code with AI Assistants?" *UC Berkeley Security Seminar* , *May 2023*
17. "Assistive Teaching of Motor Control Tasks to Humans" *Stanford Collaborative Haptics in Robotics and Medicine Lab*, *May 2023*.
18. "Assistive Teaching of Motor Control Tasks to Humans" *Simons Institute Workshop on AI & Humanity* , *July 2022*.
19. "Assistive Teaching of Motor Control Tasks to Humans" *SystemX Alliance Fall Conference* , *November 2022*.
20. "Mathematical Notions vs. Human Perception of Fairness: A Descriptive Approach to Fairness for Machine Learning" *Oxford University Algorithms at Work Reading Group*, *May 2021*.
21. "Mathematical Notions vs. Human Perception of Fairness: A Descriptive Approach to Fairness for Machine Learning" *Microsoft Research AI & Society Reading Group*, *October 2019*.
22. "Fairness & Robustness with Missing Information" *Stanford Causality & Cognition Lab*, *December 2020*.