

Megha Srivastava

megha@cs.stanford.edu • www.cs.stanford.edu/~megha
735 Campus Dr., Stanford, CA 94305

Education

Stanford University • Stanford, CA present
PhD in Computer Science, co-advised by Dorsa Sadigh & Dan Boneh
Supported by the NSF Graduate Research Fellowship (2018 – 2023) and IBM PhD Fellowship (2023-2024)
Current Research Interests: AI and education, human-AI interaction, machine learning safety/security

Massachusetts Institute of Technology • Cambridge, MA Fall 2023 – Winter 2024
Visiting PhD student at the Computer Science & Artificial Intelligence Laboratory, hosted by Jacob Andreas

Stanford University • Stanford, CA Fall 2014 – Spring 2019
BS in Computer Science with Honors and Distinction, Minor in Creative Writing
MS in Computer Science (Artificial Intelligence), advised by Percy Liang & Tatsunori Hashimoto
Ben Wegbreit Prize for Best Undergraduate Honors Thesis in Computer Science
Bing Overseas Studies Program at Oxford University (Logic & Computability) in Spring 2017

Experience

Member of Technical Staff (Contractor) – Anthropic present
San Francisco, CA
• Support research projects within the Alignment team.

AI Resident – Microsoft Research Fall 2019 – Fall 2020
Redmond, WA
• Quantifying effects of nondeterminism in neural network training (Mentors: Besmira Nushi, Ece Kamar, Eric Horvitz).

Research Intern – Google Research Summer 2019
Los Angeles, CA
• Improving robustness of image understanding models with the Research & Machine Intelligence team.

Research Intern – ETH Zürich Learning & Adaptive Systems Group Summer 2018
Zürich, Switzerland
• Active learning and adaptive hypothesis testing (Mentors: Andreas Krause, Hoda Heidari)
• Supported by the ETH Zürich Student Summer Research Fellowship.

Research Assistant – Vision & Perception Neuroscience Group Summer 2016 – Spring 2017
Stanford, CA
• Generalization and perceptual invariances in both human and artificial vision models (Mentor: Kalanit Grill-Spector)
• Supported by the Bio-X Undergraduate Research Award.

Awards & Fellowships

- Human Robot Interaction (HRI) Pioneers, 2025
- IBM PhD Fellowship, 2023
- American Association for the Advancement of Science (AAAS) Mass Media Fellowship Finalist, 2019
- International Conference on Machine Learning (ICML) Best Paper Runner-Up Award, 2018
- National Science Foundation Graduate Research Fellowship, 2018
- Ben Wegbreit Prize for Best Undergraduate Honors Thesis in Computer Science (Stanford), 2018
- Phi Beta Kappa, 2018
- Bio-X Research Award (Stanford), 2016
- Lunsford Award for Oral Presentation of Research Finalist (Stanford), 2016
- Intel Science Talent Search Semifinalist (High School), 2013

Activities & Service

- Science Small Groups Mentor (2024)
- Course Assistant: CS 255 (Introduction to Cryptography, 2025), CS 221 (Introduction to Artificial Intelligence, 2018)
- Stanford AI Lab Blog Editor, 2020-present
- Organizer for Women in Machine Learning (WiML) 2023, co-located with NeurIPS 2023
Student Program & Travel Funding Chair
- Reviewer for NeurIPS (2020 - curr.), ICLR (2021 - curr.), ICML (2021 - curr.), CoRL 2023, L4DC 2024, CHI 2025
- Stanford Women in Computer Science (Vice President), 2017 – 2018
- Photography showcased at Stanford Art of Science Exhibition, Frost Music Festival, dnjGallery at Bergamot Station, and Stanford UNICEF Arts Gallery

contd.

Preprints

Alexis Ross*, **Megha Srivastava***, Jeremiah Blanchard, Jacob Andreas. "Modeling Student Learning with 3.8 Million Program Traces", arxiv 2025

Selected Publications (See Google Scholar for full list, * denotes equal contribution)

Megha Srivastava*, Reihaneh Iranmanesh*, Yuchen Cui, Deepak Gopinath, Emily Sumner, Andrew Silva, Laporsha Dees, Guy Rosman, Dorsa Sadigh. "Shared Autonomy for Proximal Teaching", ACM/IEEE International Conference on Human-Robot Interaction (HRI), 2025

Megha Srivastava, Simran Arora, Dan Boneh. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism", Advances in Neural Information Processing Systems 37 (NeurIPS), 2024.

Neil Perry*, **Megha Srivastava***, Deepak Kumar, Dan Boneh. "Do Users Write More Insecure Code with AI Assistants?" ACM Conference on Computer and Communications Security (CCS), 2023 mlsec.org **Top-100 Computer Security Papers** 

Megha Srivastava, Erdem Biyik, Suvir Mirchandani, Noah Goodman, Dorsa Sadigh. "Assistive Teaching of Motor Control Tasks to Humans" Advances in Neural Information Processing Systems 35 (NeurIPS), 2022.

Megha Srivastava and Noah Goodman. "Question Generation for Adaptive Education" Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL), 2021.

Megha Srivastava, Tatsunori Hashimoto, Percy Liang. "Robustness to Spurious Correlations via Human Annotations" Proceedings of the 37th International Conference on Machine Learning (ICML), 2020.

Tatsunori Hashimoto, **Megha Srivastava**, Hongseok Namkoong, Percy Liang. "Fairness Without Demographics in Repeated Loss Minimization." Proceedings of the 35th International Conference on Machine Learning (ICML), 2018. **Best Paper Runner-Up Award**

Invited Talks

1. "Security of Modern Machine Learning" *Accenture Responsible AI Course*, September 2025
2. "Security of Modern Machine Learning" *AI4All*, June 2025
3. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *Google Privacy in ML Seminar*, February 2025
4. "Security of Modern Machine Learning" *DaVita x Stanford Human-Centered AI*, February 2025
5. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *AI and Cryptography Workshop, Joint Mathematics Meeting*, January 2025
6. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *University of Washington Security Seminar*, January 2025
7. "New Challenges of Trust with Large-Scale AI Systems" *University of Chicago*, November 2024.
8. Beyond Instruction Following: Language and Human-Robot Interaction *Stanford University (Guest Lecture for CS 329X: Human-Centered NLP)*, October 2024
9. "Security of Modern Machine Learning" *Stanford Cybersecurity & Privacy Festival*, October 2024
10. "Policy Learning with a Language Bottleneck" **Spotlight Talk** at *Training Agents with Foundation Models Workshop (RLC 2024)*, August 2024.
11. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *University of Toronto*, July 2024
12. "Security of Modern Machine Learning" *Accenture Responsible AI Course*, July 2024
13. "Optimistic Verifiable Training by Controlling Hardware Nondeterminism" *Stanford Security Forum*, April 2024
14. "Challenges in Human-AI Interaction for Information-Seeking Tasks" *Rising Stars in Machine Learning Workshop*, November 2023
15. "Do Users Write More Insecure Code with AI Assistants?" *University of Maryland College Park (Guest Lecture for Large Language Models, Security, and Privacy seminar)*, October 2023.
16. "Robustness to Spurious Correlations via Human Annotations" *Two Sigma PhD Symposium*, June 2023.

Invited Talks (cont.)

17. "Do Users Write More Insecure Code with AI Assistants?" *UC Berkeley Security Seminar*, May 2023
18. "Assistive Teaching of Motor Control Tasks to Humans" *Stanford Collaborative Haptics in Robotics and Medicine Lab*, May 2023.
19. "Assistive Teaching of Motor Control Tasks to Humans" *Simons Institute Workshop on AI & Humanity*, July 2022.
20. "Assistive Teaching of Motor Control Tasks to Humans" *SystemX Alliance Fall Conference*, November 2022.
21. "Mathematical Notions vs. Human Perception of Fairness: A Descriptive Approach to Fairness for Machine Learning" *Oxford University Algorithms at Work Reading Group*, May 2021.
22. "Mathematical Notions vs. Human Perception of Fairness: A Descriptive Approach to Fairness for Machine Learning" *Microsoft Research AI & Society Reading Group*, October 2019.
23. "Fairness & Robustness with Missing Information" *Stanford Causality & Cognition Lab*, December 2020.